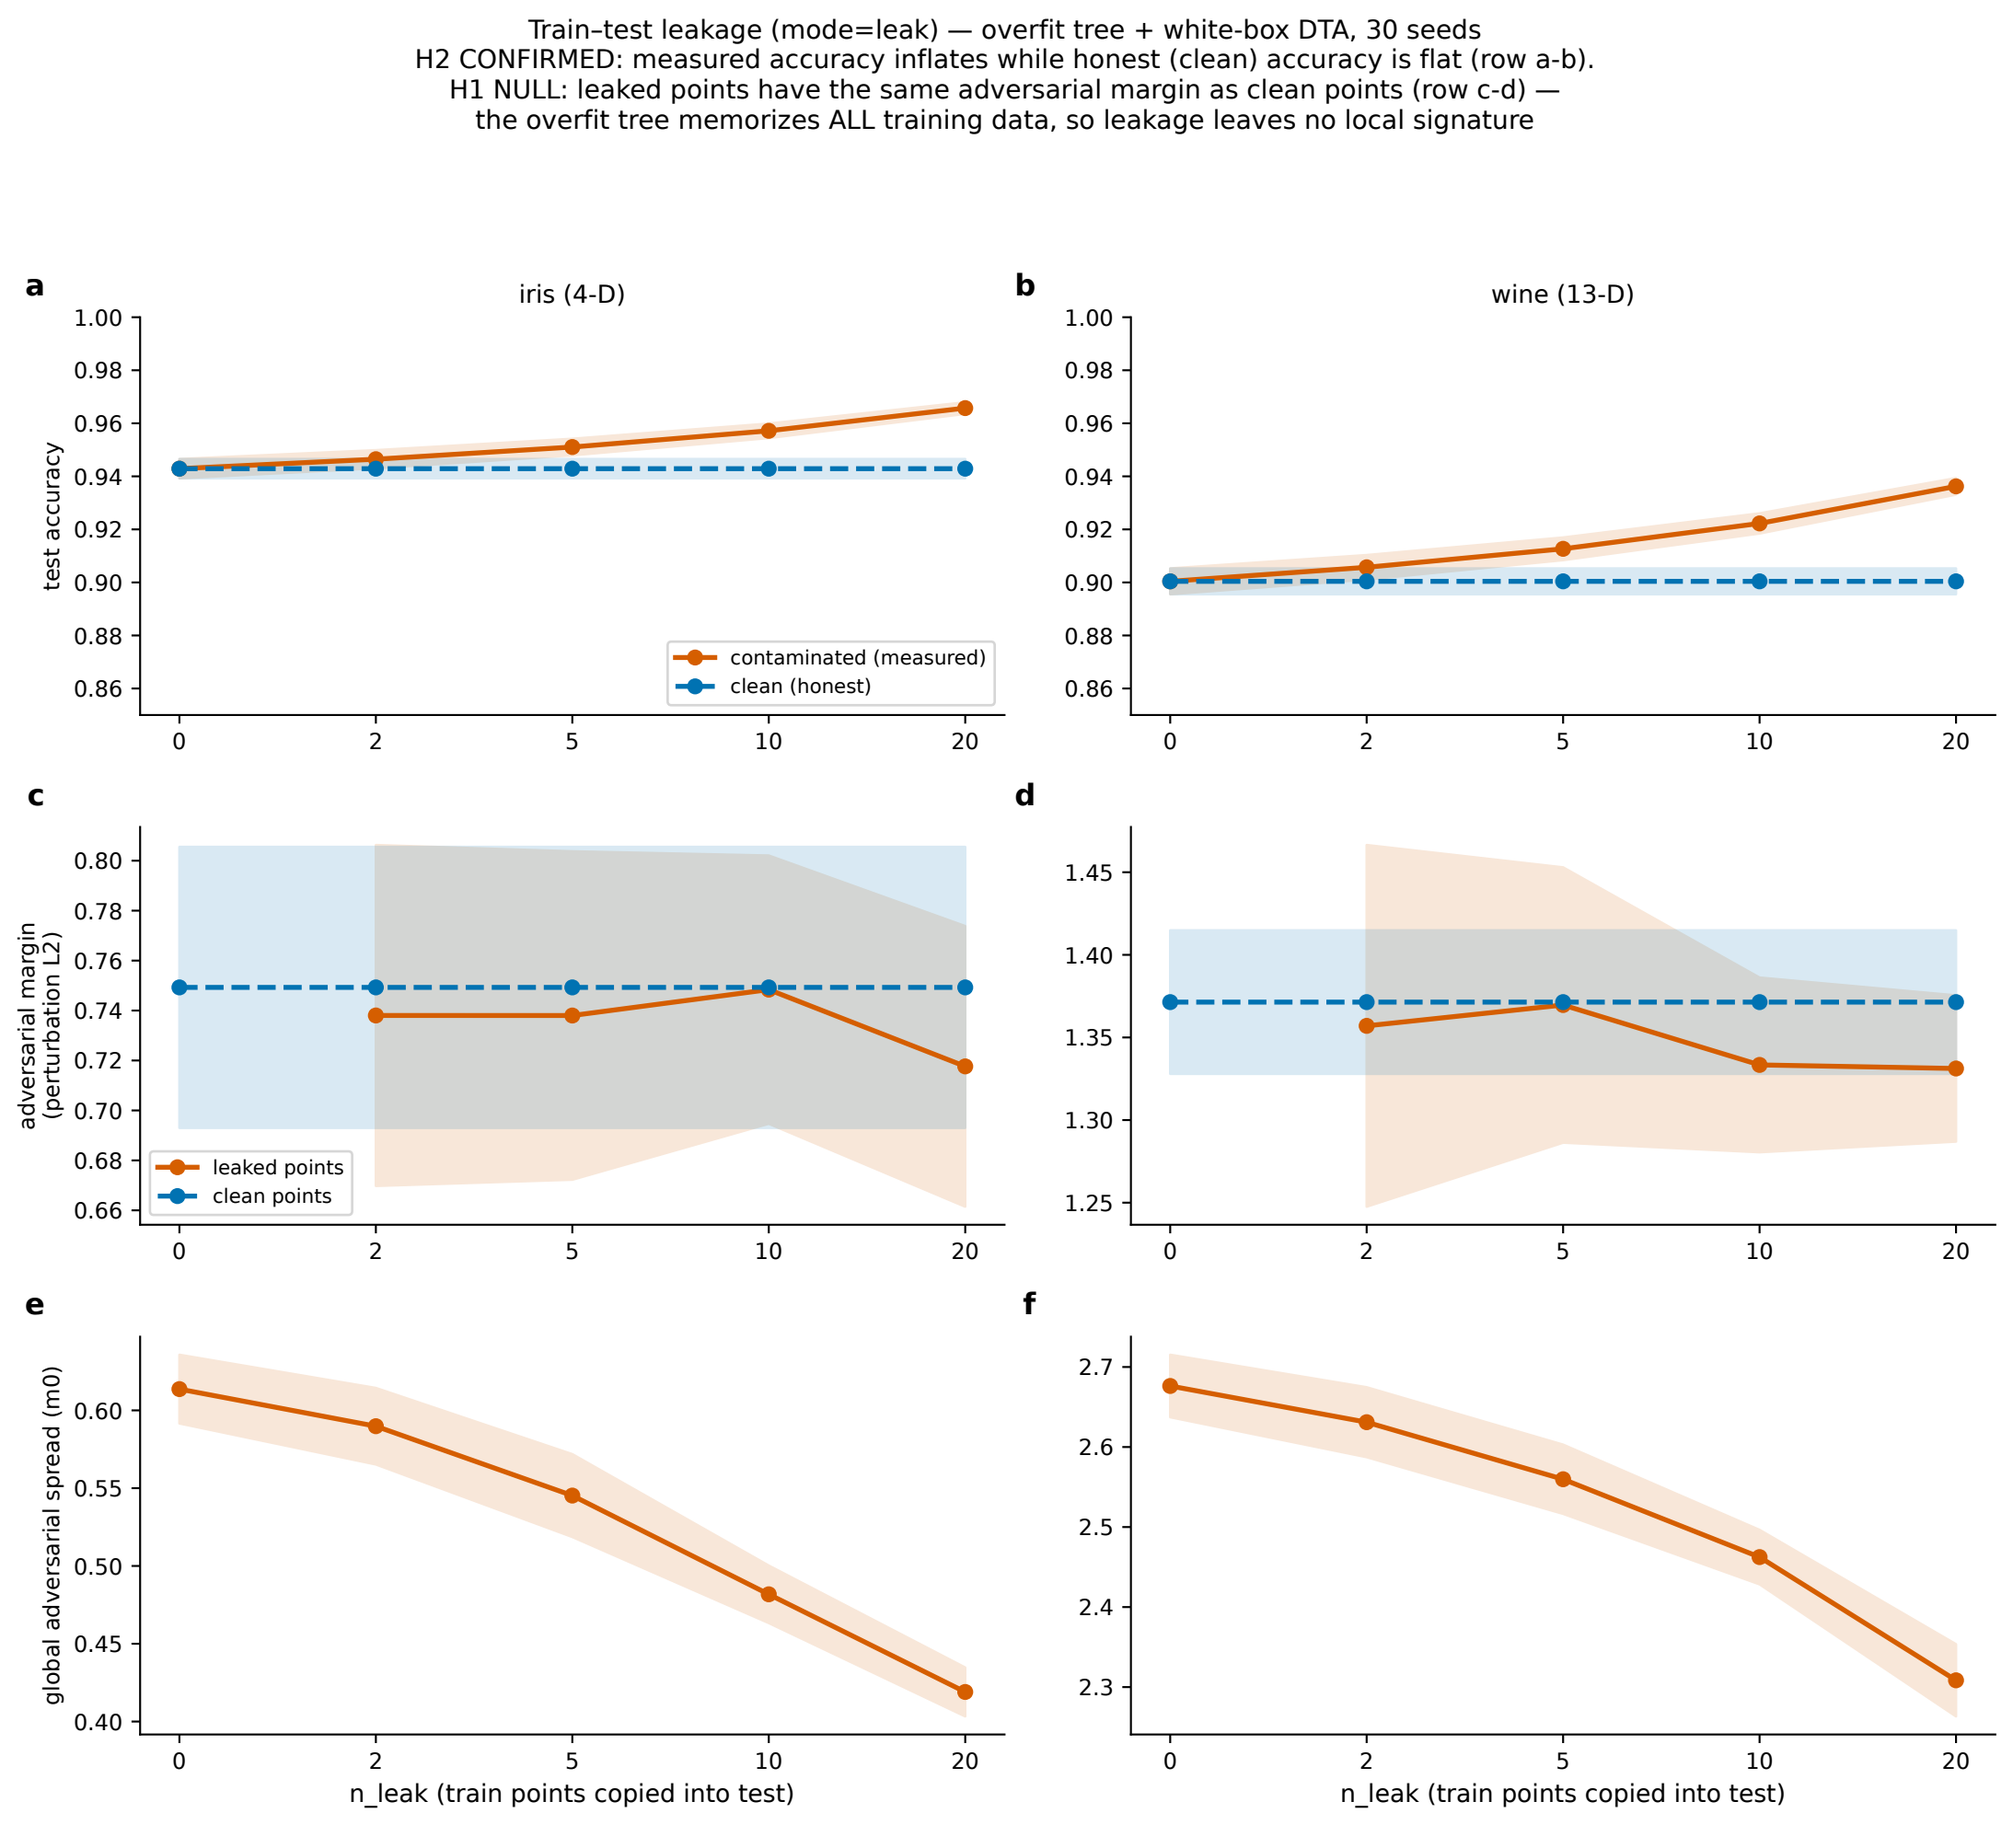
Train–test leakage (mode=leak) — overfit tree + white-box DTA, 30 seeds
H2 CONFIRMED: measured accuracy inflates while honest (clean) accuracy is flat (row a-b).
H1 NULL: leaked points have the same adversarial margin as clean points (row c-d) —
the overfit tree memorizes ALL training data, so leakage leaves no local signature
a
iris (4-D)
b
wine (13-D)
test accuracy
contaminated (measured)
clean (honest)
c
d
adversarial margin
(perturbation L2)
leaked points
clean points
e
f
global adversarial spread (m0)
n_leak (train points copied into test)
n_leak (train points copied into test)